

Responsible Leadership for Societal-Scale AI: A Human-Centered Socio-Technical Approach to Ethical AI Deployment in Healthcare

Abhijit Khompi

Doctor of Business Administration Program, Edward S. Ageno School of Business, Golden Gate University, San Francisco, CA, USA

Abstract

Healthcare organizations are deploying artificial intelligence at societal scale across clinical decision support, diagnostics, operations, and population health yet an estimated 70–90% of digital health initiatives fail to deliver expected value (Topol, 2019). These failures are rarely purely technical. They reflect leadership that optimizes algorithms without jointly optimizing the human systems, trust, professional identity, ethical norms, and workforce wellbeing, on which responsible AI depends. In resource-constrained settings, frugal and societally accountable AI requires leaders who steward technology without sacrificing human-centered care.

This research paper extends Socio-Technical Systems (STS) theory [31, 32] as a framework for responsible, human-centered AI deployment. STS holds that organizational effectiveness emerges only when social and technical subsystems are jointly optimized. The study examines leadership effectiveness in AI-integrated healthcare by operationalizing the social subsystem through ethical and adaptive leadership [34, 35] and the technical subsystem through digital leadership and AI technology acceptance [36, 37]. The purpose is to determine how leaders who balance technological advancement with human values enable responsible AI governance and improved outcomes.

A sequential explanatory mixed-methods design will be employed. The quantitative phase surveys approximately 150 healthcare leaders using structural equation modeling; the qualitative phase conducts semi-structured interviews with 12–15 leaders to interpret how joint optimization is enacted in practice. Expected contributions include a validated STS-based measurement model for AI-era leadership, a practical framework for responsible AI implementation under healthcare resource constraints, and evidence-based insights on balancing technological capability with human values. Societal impact is framed as safer patient care, more equitable algorithmic outcomes, and sustainable workforce conditions in AI-enabled health systems.

Keywords: *responsible AI, human-centered AI, socio-technical systems, ethical leadership, healthcare AI, societal-scale AI*

1 Introduction

Like many industries Artificial intelligence is also reshaping healthcare at a substantial pace and scale unlike anything reported in the history of clinical medicine. From algorithmic triage in emergency departments to machine learning models that flag sepsis risk before clinical signs manifest, and from diagnostic imaging systems that rival board-certified radiologists to natural language processing tools that surface actionable insights from unstructured clinical notes, AI has moved from pilot programs into production clinical workflows in hospitals and health systems across multiple countries [1, 2]. According to industry analyses, the healthcare AI market is projected to exceed USD 67 billion by 2027, with applications spanning clinical decision support, population health management, administrative automation, and precision medicine [3]. In settings constrained by workforce shortages, rising acuity, and fiscal pressure, the promise of AI as a force multiplier for human capability is both compelling and urgent.

Yet this promise is being systematically squandered. Across the peer-reviewed literature and industry intelligence, a sobering consensus has emerged: an estimated 70-90% of digital health initiatives fail to

deliver their anticipated value [1, 4]. These failures are not predominantly failures of technology. Clinical AI tools that perform impressively in controlled validation studies frequently fail to integrate into clinical workflows, fail to earn the trust of clinicians, or fail to reduce rather than exacerbate existing health disparities when deployed in real-world settings [1, 5]. The algorithm that Obermeyer and colleagues exposed in a landmark 2019 Science study illustrates the stakes: a widely deployed, commercially successful tool used to manage care for millions of patients systematically underestimated the health needs of Black patients relative to white patients with equivalent illness burden, producing racially biased allocation of follow-up care [1]. The algorithm did not fail technically; it succeeded at the objective it was given. What failed was the governance, the leadership judgment about what objectives to encode, and the oversight that should have caught the human consequences before deployment at scale.

This paper examines how leadership practices not technology alone determine whether AI deployments in healthcare produce the patient safety, equity, and workforce outcomes they are designed to achieve. When organizations invest in AI technology without equivalent investment in the human and organizational conditions on which that technology depends trust, professional identity, ethical culture, workforce preparation, and adaptive capacity they are building on sand. Responsible AI is not primarily an engineering problem; it is a leadership and governance problem. Leaders who optimize algorithms without jointly optimizing the human systems surrounding them produce AI that is technically capable but organizationally and ethically fragile [6, 7]. In resource-constrained settings, where frugal AI is not a philosophical preference but an operational necessity, this leadership gap is especially consequential: there is little margin for failed deployments, eroded clinician trust, or inequitable algorithmic outcomes.

The theoretical framework most suited to this problem is Socio-Technical Systems (STS) theory, which holds that organizational effectiveness emerges only when social and technical subsystems are jointly designed and optimized [8, 9]. Originally developed by Trist and Bamforth in the context of coal mining and extended by Cherns into a set of design principles, STS theory has proved remarkably generative across industries and technological transitions. Yet its application to AI-era leadership in healthcare remains theoretically underdeveloped. This paper extends STS theory to address the specific challenges of responsible AI deployment in healthcare, operationalizing the social subsystem through ethical and adaptive leadership and the technical subsystem through digital leadership competencies and AI technology acceptance.

1.1 Research Objectives and Contributions

The central research question I formulated to guide this work is: How do healthcare leaders who jointly optimize social and technical subsystems enable responsible AI governance and improved organizational outcomes in AI-integrated clinical environments? Three interrelated objectives follow from this question. First, to extend STS theory into a leadership-mediated model of responsible AI governance that is theoretically coherent and empirically testable. Second, to operationalize the social and technical subsystems of AI-integrated healthcare through validated leadership constructs ethical leadership, adaptive leadership, digital leadership, and AI technology acceptance facilitation and to specify the joint optimization mechanism that links these constructs to governance quality. Third, to generate a practical framework that healthcare leaders can apply under resource constraints.

The paper makes four primary contributions. (1) Theoretical: it advances STS theory by specifying leadership as the primary mechanism of joint optimization in AI-era healthcare organizations, filling a

gap between classical STS design principles and contemporary responsible AI governance. (2) Conceptual: it introduces the ‘Khompi’s Responsible AI Leadership Framework’ (RALF), integrating four leadership construct clusters as determinants of joint optimization, responsible AI governance, and organizational AI outcomes. (3) Methodological: it specifies a sequential explanatory mixed-methods design combining structural equation modeling of approximately 150 healthcare leaders with interpretive interviews of 12–15 leaders. (4) Practical: it offers actionable guidance for healthcare executives, clinical leaders, and policymakers seeking to deploy AI in ways that advance technological capability without sacrificing clinician agency, patient trust, or workforce wellbeing.

The significance of this research extends beyond academic theory-building. Healthcare is a critical societal vertical: the quality of leadership in AI-integrated health systems shapes care quality, clinical equity, and workforce wellbeing at population scale. As healthcare organizations worldwide confront the dual imperatives of technological modernization and resource constraint, leaders need validated frameworks for deploying AI responsibly rather than merely rapidly. The remainder of this paper proceeds as follows. Section 2 synthesizes STS theory, leadership research, AI adoption scholarship, and healthcare AI governance debates. Section 3 describes the sequential explanatory mixed-methods design. Section 4 presents the Responsible AI Leadership Framework, its theoretical propositions, structural equations, and an illustrative clinical application. Section 5 discusses implications for practice & policy. Section 6 concludes and outlines future scope.

2 Literature Review

2.1 Socio-Technical Systems Theory: Foundations and Evolution

Socio-Technical Systems (STS) theory originated in Trist and Bamforth’s (1951) field study of mechanization in English coal mines. In 1951, Eric Trist and Ken Bamforth of the Tavistock Institute published a study of the longwall method of coal-getting that would fundamentally challenge prevailing assumptions about industrial organization [8]. They observed that the introduction of mechanized longwall technology had been accompanied by a dramatic deterioration in productivity, worker wellbeing, and organizational morale despite the technology’s self-evident mechanical superiority over the prior hand-got method. The reason, they argued, was that technology had been introduced as though it existed in a social vacuum. The new machinery destroyed the small, self-regulating work groups and the informal social systems that had sustained the prior method, replacing human interdependence and skill with fragmented, surveillance-intensive task structures. Technology had been optimized; the human system had been ignored; and the resulting misfit between technical capability and social structure produced organizational failure [8].

Cherns subsequently codified STS theory into nine design principles, of which joint optimization is foundational: the social and technical subsystems must be designed in relation to each other and in relation to the demands of the environment, not independently [9]. Other Cherns principles—minimal critical specification (design only what is necessary), boundary management, multi-functionality, and information flow describe organizational conditions that sustain joint optimization over time. Subsequent scholars extended STS to information systems, electronic health records, and complex adaptive healthcare systems. Sittig and Singh’s sociotechnical model for health information technology identifies eight interacting dimensions hardware, software, clinical content, people, workflow, internal organizational

policies, external environment, and measurement that must be jointly addressed for HIT implementations to succeed [10]. Yet existing healthcare STS models have not been systematically applied to the leadership-mediated governance of clinical AI systems. This paper addresses that gap by positioning leadership as the active agent of joint optimization rather than treating joint optimization as a passive design property.

2.2 Leadership in Digital Transformation: Ethical, Adaptive, and Digital Leadership

Three leadership theories are particularly relevant to the social subsystem in AI-integrated healthcare: ethical leadership, adaptive leadership, and digital leadership. Brown, Treviño, and Harrison's 2005 social learning theory of ethical leadership defines ethical leaders as individuals who demonstrate normatively appropriate conduct through personal actions and interpersonal relationships, and who promote such conduct among followers through two-way communication, reinforcement, and decision-making [5]. Ethical leadership is not merely rule-following; it is actively modeling moral behavior and creating organizational conditions in which ethical conduct by others is expected, noticed, and rewarded. In AI-integrated healthcare, ethical leadership operationalizes the human dimension of responsible AI: leaders who communicate the moral stakes of algorithmic decision-making, who protect clinicians who raise concerns about AI recommendations, and who create psychological safety for questioning data-driven outputs exercise ethical leadership [5, 11].

Heifetz, Grashow, and Linsky's adaptive leadership theory distinguishes between technical problems for which expertise and established solutions suffice and adaptive challenges, which require changes in values, beliefs, roles, and working relationships that cannot be resolved through existing knowledge [6]. AI deployment in healthcare is paradigmatically an adaptive challenge. It does not merely require new technical skills; it requires clinical professionals to renegotiate their relationship with expertise, authority, and professional identity when algorithms become active participants in clinical reasoning. Leaders who treat AI adoption as a technical problem to be managed through training programs and implementation timelines miss the adaptive dimension entirely. Adaptive leaders, by contrast, create holding environments in which conflicting values can surface, work productively, and generate new organizational norms [6].

Digital leadership, as operationalized by Zeike and colleagues, refers to the competencies that enable leaders to navigate digital transformation while sustaining the psychological wellbeing of their teams [7]. These competencies include technology orientation (understanding and communicating the strategic significance of digital tools), change facilitation (supporting teams through digital transitions), and digital vision (articulating how technology serves human and organizational purpose). Zeike et al. found significant positive associations between digital leadership and employee psychological wellbeing in digitally transforming organizations, suggesting that how leaders frame and mediate technological change matters as much as what technological change entails [7]. For healthcare AI, digital leadership provides the bridge between technical capability and human meaning: the leader who can articulate why an AI diagnostic tool serves the mission of compassionate care reduces the experienced conflict between algorithmic efficiency and professional vocation.

2.3 Technology Acceptance in Healthcare: Adoption, Trust, and Explainability

The technical subsystem in AI-integrated healthcare includes not only the algorithms, infrastructure, and data architectures on which AI depends, but also the human-technology interface through which

clinicians adopt, adapt to, and exercise judgment about AI recommendations. Venkatesh, Thong, and Xu's Unified Theory of Acceptance and Use of Technology (UTAUT2) identifies performance expectancy, effort expectancy, social influence, facilitating conditions, hedonic motivation, habit, and price value as key determinants of technology adoption [12]. In healthcare contexts, performance expectancy the degree to which a clinician believes an AI tool will improve clinical performance is the strongest predictor of adoption intention, followed closely by facilitating conditions, which include organizational support, training, and technical infrastructure [12]. These determinants are inherently leadership-mediated: performance expectancy depends partly on how leaders communicate what an AI tool is designed to do and how well; facilitating conditions depend directly on resource allocation and infrastructure investment decisions that leaders control.

A critical dimension of technology acceptance in healthcare AI is trust and explainability. Clinicians frequently report reluctance to act on AI recommendations when they cannot understand the basis for those recommendations a phenomenon sometimes called algorithm aversion and sometimes, more charitably, appropriate epistemic caution [13]. The opacity of many high-performing AI models (particularly deep learning systems) creates a structural tension: models optimized for predictive accuracy are often less interpretable, while interpretable models may sacrifice accuracy. Leaders who champion explainable AI tools designed to make their reasoning transparent to users operationalize the ethical principle that clinicians must retain informed decision-making authority over AI-assisted recommendations. They simultaneously address both the ethical and acceptance dimensions of responsible AI [13, 14].

2.4 Healthcare AI: Promise, Peril, and the Governance Imperative

The clinical promise of AI in healthcare is substantial and growing. Machine learning models have demonstrated diagnostic accuracy at or above expert clinician level in radiology, dermatology, ophthalmology, and pathology [15]. Natural language processing has enabled extraction of clinically meaningful information from electronic health records at scale, transforming unstructured narrative into structured intelligence. Predictive analytics tools have demonstrated capacity to identify patients at high risk of deterioration, readmission, or sepsis before clinical recognition, enabling earlier intervention and resource allocation [16]. At the population level, AI offers tools for disease surveillance, genomic medicine, and the optimization of public health resource deployment. The technological capability is real.

So is the peril. The Obermeyer et al. finding of systematic racial bias in a widely deployed commercial algorithm is not an anomaly but a cautionary archetype [1]. Similarly, Larrazabal et al. demonstrated that chest X-ray diagnostic AI models trained predominantly on data from one sex exhibited significant performance degradation when applied to the other sex, a finding that underscores the need for the equity auditing Khompi's RALF requires before deployment [38]. The same structural conditions optimization for proxies rather than health needs, training on historically biased data, deployment without adequate equity auditing characterize dozens of clinical AI tools currently in use. Beyond algorithmic bias, healthcare AI raises concerns about erosion of clinical judgment, data privacy and sovereignty, and the displacement of healthcare workers without adequate transition support [24, 29]. At the governance level, the World Health Organization's 2021 Ethics and Governance of AI for Health guidance identifies six core principles for responsible health AI: protecting human autonomy, promoting human wellbeing and safety, ensuring transparency, fostering accountability, ensuring equity and inclusiveness, and promoting AI that is responsive and sustainable [24]. These principles are not self-executing; they require leaders

who embody and institutionalize them through governance structures, incentive systems, and cultural norms [25, 27]. Table 1 maps the WHO principles to the Khompi's RALF leadership constructs developed in this paper, illustrating how ethical, adaptive, and digital leadership jointly operationalize global governance norms at the organizational level.

Table 1. Mapping WHO Ethics and Governance of AI for Health principles to Khompi's RALF leadership constructs.

WHO Principle	Primary Khompi's RALF Construct(s)	Leadership Enactment in Healthcare AI
Protect human autonomy	Ethical + Adaptive leadership	Preserve clinician agency; treat AI as input to judgment, not substitute
Promote wellbeing and safety	Joint optimization + Digital leadership	Pace deployment to readiness; monitor patient safety alongside accuracy
Ensure transparency	Digital leadership + ATA facilitation	Champion explainable AI; communicate performance data openly
Foster accountability	Ethical leadership	Reward appropriate override; hold abdication of judgment accountable
Ensure equity and inclusiveness	Ethical leadership + Joint optimization	Mandate equity audits before and after deployment
Responsive and sustainable AI	Adaptive + Digital leadership	Build adaptive capacity; invest in sustainable workforce conditions

2.5 Synthesis: The Leadership Gap in Healthcare AI Governance

A review of the extant literature reveals a structural gap. STS theory provides a powerful framework for understanding why organizational effectiveness requires joint optimization of social and technical subsystems. Leadership theory provides validated constructs for the ethical, adaptive, and digital competencies that constitute effective leadership in digital transformation. Technology acceptance theory identifies the adoption-mediating roles of performance expectancy, social influence, and facilitating conditions. Healthcare AI governance frameworks articulate the principles to which responsible deployment must conform. What is missing is an integrated, leadership-centered model that connects these streams and provides healthcare leaders with actionable guidance for governing AI deployment in ways that simultaneously advance technological capability and protect human values. This paper proposes such a model, grounded in STS theory and designed for empirical validation through a rigorous mixed-methods research design.

This gap persists for an identifiable structural reason: the academic disciplines that study AI governance (computer science, information systems) and those that study healthcare leadership (organizational behavior, health services research) operate in largely separate publication ecosystems, attend different conferences, and train doctoral students in different methodological traditions. Khompi's RALF is designed to bridge this disciplinary divide.

3 Materials and Methodology

3.1 Research Design

A sequential explanatory mixed-methods design is employed, in which a quantitative phase is followed by a qualitative phase that interprets and extends the quantitative findings [17]. This design is methodologically appropriate for three reasons. First, the research questions require both measurement and meaning: the quantitative phase tests structural relationships among leadership constructs and outcomes, while the qualitative phase illuminates how those relationships are enacted in the specific organizational and cultural contexts of AI-integrated healthcare. Second, the explanatory sequence allows qualitative sampling to be purposively directed by quantitative findings, maximizing the interpretive yield of the qualitative phase. Third, mixed methods triangulation strengthens validity by enabling cross-method convergence and divergence to be identified and interrogated. Figure 1 (presented in Section 4) depicts the conceptual model that the design is intended to test; the methodological architecture that generates evidence for that model is specified below.

3.2 Quantitative Phase: Structural Equation Modeling Survey

The quantitative phase surveys approximately 150 healthcare leaders using a structured instrument administered via validated survey platforms. Participants are recruited from hospitals, integrated health systems, and large ambulatory care networks in which AI-enabled clinical decision support tools are currently deployed. Inclusion criteria require active leadership roles (supervisor level or above) with direct responsibility for AI implementation, governance, or clinical workflow design. The target sample of 150 is consistent with recommended minimum thresholds for structural equation modeling (SEM) with the proposed model complexity [18]. Sample size calculations based on the anticipated number of latent constructs (seven) and indicators per construct (three to five) confirm adequacy at a target power of 0.80 and effect size f^2 of 0.15.

Validated measurement scales are adapted from established instruments, as summarized in Table 2. Ethical leadership is measured using the 10-item Ethical Leadership Scale (ELS) developed by Brown et al. [5]. Adaptive leadership is assessed through a 12-item scale adapted from the Adaptive Leadership Questionnaire, capturing behaviors of observing, interpreting, and intervening on adaptive challenges [6]. Digital leadership is measured using Zeike et al.'s eight-item Digital Leadership Scale [7]. AI technology acceptance is operationalized through performance expectancy, effort expectancy, and social influence subscales of UTAUT2 [12], adapted for clinical AI contexts. Outcome variables include responsible AI governance (measured through items assessing fairness auditing practices, transparency mechanisms, and accountability structures) and organizational AI outcomes (a composite of patient safety indicators, clinician acceptance rates, and operational efficiency gains attributable to AI deployment). A five-point Likert response format is used throughout.

Table 2. Measurement constructs, source instruments, and indicator counts for the quantitative phase.

Construct	Source Instrument	Items	Subsystem / Role
Ethical Leadership (EL)	ELS (Brown et al., 2005)	10	Social subsystem
Adaptive Leadership (AL)	Adaptive Leadership Questionnaire	12	Social subsystem
Digital Leadership (DL)	Digital Leadership Scale (Zeike et al., 2019)	8	Technical subsystem

AI Tech. Acceptance Facilitation (ATAF)	UTAUT2 subscales (Venkatesh et al., 2012)	9	Technical subsystem
Joint Optimization (JO)	Author-developed STS composite	8	Core mechanism
Responsible AI Governance (RAG)	Author-developed governance scale	7	Mediating outcome
Organizational AI Outcomes (AO)	Composite safety / equity / wellbeing	6	Distal outcome

Partial Least Squares Structural Equation Modeling (PLS-SEM) is employed for analysis. PLS-SEM is appropriate for this study given the exploratory-to-confirmatory nature of the model, the expected non-normal distribution of some constructs, and the emphasis on predictive relevance in addition to theory testing [18]. Measurement model assessment evaluates indicator reliability (factor loadings ≥ 0.70), internal consistency (Cronbach's α and composite reliability ≥ 0.80), convergent validity (average variance extracted ≥ 0.50), and discriminant validity (HTMT criterion < 0.90). Structural model assessment tests path coefficients and their significance via bootstrapping (5,000 iterations), R-squared values for endogenous constructs, and effect sizes (f^2). Common method bias is assessed through Harman's single-factor test and procedural remedies including temporal and spatial separation of predictor and outcome items.

3.3 Qualitative Phase: Semi-Structured Interviews

The qualitative phase conducts semi-structured interviews with 12–15 healthcare leaders selected through purposive sampling informed by quantitative results. Participants are selected to represent variation in quantitative response patterns, organizational context (academic medical centers, community hospitals, integrated health systems), and geographic diversity. This sample size is consistent with established recommendations for theoretical saturation in qualitative leadership research [19]. Interviews are conducted via secure video conferencing, recorded with participant consent, and transcribed verbatim. Each interview lasts approximately 60 minutes.

Interview protocol development follows the exploratory-confirmatory mixed-methods convention: core questions address how leaders enact joint optimization in practice, how they navigate conflicts between technological imperatives and human values, what governance structures they have established for AI oversight, and how they sustain clinician trust and agency in AI-integrated workflows. Probes are designed to surface specific incidents, decisions, and turning points rather than general attitudes, maximizing the explanatory depth of the qualitative data. Interview data are analyzed through template analysis [19], beginning with an a priori template derived from STS and leadership theory and refined inductively through iterative engagement with the data. A second analyst independently codes 25% of the transcripts to assess inter-rater reliability; discrepancies are resolved through discussion until consensus is reached. NVivo software supports data management and audit trail.

3.4 Ethical Considerations

I have already received IRB approval on study protocol from Golden Gate University prior to data collection. Participation is voluntary; written informed consent is obtained from all participants. Quantitative data are anonymized at the point of collection; qualitative transcripts are de-identified prior to analysis. Data are stored in encrypted cloud environments accessible only to the research team. Findings are reported in aggregate with no individual organizations or leaders identifiable from reported data [21]. Given that participants are senior leaders in healthcare organizations, risks of participatory

harm are minimal; the primary ethical obligation is accurate representation of findings without distortion that could misrepresent individual or organizational practices.

4 Proposed Model: The Responsible AI Leadership Framework

4.1 Framework Overview

I structured Khompi's Responsible AI Leadership Framework (RALF) around the joint optimization mechanism, rather than treating ethical and digital leadership as independent predictors, because my review of the STS literature convinced me that Cherns' (1976) foundational insight has been systematically underappreciated in the AI governance discourse: it is the *interaction* between social and technical subsystems, not their independent quality, that determines organizational effectiveness is as an STS-grounded, leadership-mediated model for governing AI deployment in healthcare. Khompi's RALF integrates four leadership construct clusters ethical leadership, adaptive leadership, digital leadership, and facilitation of AI technology acceptance as determinants of the joint optimization of social and technical subsystems, which in turn predicts responsible AI governance quality and, ultimately, improved organizational AI outcomes. I chose to include both ethical and adaptive leadership as separate social subsystem constructs, rather than collapsing them, because they address fundamentally different challenges: ethical leadership addresses the *normative* dimension (what ought we to do?), while adaptive leadership addresses the *identity* dimension (who are we becoming when AI changes our work?). The model is depicted schematically in Fig. 1.




Fig. 1. The Khompi's Responsible AI Leadership Framework (Khompi's RALF): An STS-grounded model for healthcare AI governance. Social subsystem constructs (ethical leadership, adaptive leadership) and technical subsystem constructs (digital leadership, AI technology acceptance facilitation) jointly determine joint optimization quality, which predicts responsible AI governance and organizational AI outcomes. Resource constraints moderate the joint-optimization-to-governance path, reflecting frugal AI conditions.

I built Khompi's RALF's conceptual architecture on two foundational claims I have derived from STS theory. First, healthcare AI deployment involves two interdependent subsystems that must be jointly optimized: a social subsystem comprising the human, cultural, and relational dimensions of clinical work, and a technical subsystem comprising the algorithms, data infrastructures, workflow integrations, and adoption conditions on which AI depends. Second, leadership is the primary mechanism through which joint optimization is achieved or undermined: it is leaders who decide what technologies to deploy, at what pace, with what governance structures, and with what attention to the human consequences of technological change [8, 9, 20].

4.2 Social Subsystem: Ethical and Adaptive Leadership

The social subsystem in AI-integrated healthcare encompasses clinician professional identity, trust (in AI, in leadership, and in institutional intentions), workforce wellbeing, ethical culture, and the relational fabric of clinical teams. Leaders who do not attend to these dimensions when deploying AI create conditions for workforce disengagement, algorithmic over-reliance, and governance failure. Khompi's RALF identifies ethical leadership and adaptive leadership as the primary social subsystem constructs because they capture distinct but complementary dimensions of the human-centered work of responsible AI deployment.

Ethical leadership's contribution to joint optimization operates through three mechanisms. First, role modeling: when leaders visibly engage in ethical reasoning about AI questioning algorithmic recommendations, raising equity concerns, and treating patient dignity as a constraint rather than a variable they signal to their organizations that technical optimization is bounded by human values [30]. Second, transactional mechanisms: ethical leaders use reward structures and accountability systems to reinforce ethically congruent AI use, for example recognizing clinicians who appropriately override biased recommendations and holding accountable those who abdicate clinical judgment to algorithms [25]. Third, communication: ethical leaders explain the moral stakes of AI decisions in concrete human terms, reducing the psychological distance between algorithmic abstraction and patient impact.

Adaptive leadership addresses the distinctive challenge that AI deployment imposes on professional identity and established ways of working. When an AI diagnostic tool performs at expert level, it does not merely add a new capability; it raises a destabilizing question for clinicians: what is the nature and value of my expertise when an algorithm can match or exceed it? This is an adaptive challenge, not a technical one [6]. Leaders who diagnose this adaptation work accurately recognizing that resistance to AI often reflects legitimate anxiety about professional meaning rather than mere technological conservatism can facilitate productive conversations about how human expertise and algorithmic capability are complementary rather than competing. Adaptive leaders create the psychological holding environment in which clinicians can grieve the loss of familiar ways of working, experiment with new roles, and construct meaningful professional identities in AI-augmented practice.

4.3 Technical Subsystem: Digital Leadership and AI Technology Acceptance

The technical subsystem encompasses the AI systems themselves, the data and infrastructure on which they depend, the workflow integrations through which clinicians encounter AI recommendations, and the adoption conditions that determine whether technically capable tools are actually used effectively. Khompi's RALF identifies digital leadership and AI technology acceptance facilitation as the primary

technical subsystem constructs because they capture, respectively, the strategic and operational dimensions of the human-technical interface.

Digital leadership's role in the technical subsystem is strategic and enabling. Leaders with strong digital leadership competencies understand the capabilities and limitations of AI technology with sufficient depth to make sound deployment decisions: which clinical use cases are technically mature enough for deployment, what governance structures are required, what data governance protections are non-negotiable, and what implementation pace is consistent with organizational readiness. They communicate a digital vision that situates AI as a means to clinical excellence and human flourishing rather than as an end in itself, thereby aligning technical ambition with organizational purpose [7]. They also make the investment decisions in infrastructure, training, technical support, and performance monitoring that constitute the facilitating conditions without which clinician AI adoption cannot be sustained.

AI technology acceptance facilitation addresses the adoption dimension of the technical subsystem. UTAUT2 identifies multiple determinants of technology adoption intention, several of which are leadership-mediated [12]. Performance expectancy is shaped by how leaders communicate and demonstrate AI capabilities; social influence operates through leader endorsement and peer network effects that leaders can cultivate; facilitating conditions are directly determined by leader resource allocation. Leaders who proactively address adoption barriers by ensuring adequate training, by creating feedback channels for clinician concerns about AI performance, by advocating for user-centered AI design that minimizes workflow friction, and by transparently communicating performance data about deployed AI tools operationalize the technical subsystem dimension of responsible AI governance.

4.4 Joint Optimization as the Core Mechanism

Khompi's RALF's central theoretical claim is that joint optimization of social and technical subsystems is the mechanism through which responsible AI governance is achieved. Consistent with STS theory, the model does not treat social and technical optimization as additively related to outcomes; rather, it is the joint quality of their alignment that determines governance quality. Two organizations may achieve similar scores on individual ethical leadership or digital leadership dimensions yet differ dramatically in responsible AI governance outcomes if one achieves joint optimization and the other does not [8, 9].

Joint optimization in AI-integrated healthcare can be operationalized as the degree to which AI deployment decisions simultaneously advance technical performance objectives and protect or enhance social system conditions. This is measured in Khompi's RALF through a composite construct assessing: (a) the extent to which governance processes require social impact assessment alongside technical performance evaluation; (b) the degree to which AI implementation timelines are paced to organizational and clinical readiness rather than driven by technological availability alone; (c) the depth of clinician involvement in AI selection, design, and governance; and (d) the presence of post-deployment monitoring mechanisms that track human and organizational outcomes as well as technical performance indicators [20, 24]. Table 3 contrasts a technical-optimization-only approach with a Khompi's RALF-informed joint optimization approach across key deployment dimensions.

In my analysis of the healthcare AI governance literature, I observed that the term 'responsible AI' is frequently invoked without operationalization organizations claim responsible deployment without

specifying what leadership behaviors, governance structures, or measurement systems constitute responsibility in practice. Khompi's RALF's contribution is precisely this operationalization.

Table 3. Technical-optimization-only versus Khompi's RALF joint-optimization approaches to healthcare AI deployment.

Deployment Dimension	Technical-Only Approach	Khompi's RALF Joint-Optimization Approach
Algorithm selection	Maximize accuracy / AUC on validation set	Accuracy plus equity audit across demographic subgroups
Implementation pace	Driven by vendor timeline and ROI targets	Paced to clinical and organizational readiness
Clinician involvement	Training after go-live; limited design input	Co-design, governance seats, continuous feedback loops
Success metrics	Sensitivity, specificity, cost savings	Safety, equity, trust, wellbeing, and technical performance
Override culture	Discouraged as non-compliance	Protected and rewarded as clinical judgment
Post-deployment review	Technical performance dashboards	Equity, safety, workforce, and technical dashboards

4.5 Theoretical Propositions and Structural Model

From this theoretical development, I derived seven formal propositions, each designed to be testable through empirical research (Table 4). These propositions will be tested through the quantitative phase of the research design described in Section 3.

The structural relationships among Khompi's RALF constructs are conceptualized as follows:

Proposition Cluster 1 (Antecedents of Joint Optimization): Joint Optimization quality is determined by the combined influence of four leadership constructs: Ethical Leadership, Adaptive Leadership, Digital Leadership, and AI Technology Acceptance Facilitation. Each construct contributes uniquely to the joint alignment of social and technical subsystems, though their relative contributions will be empirically determined through future structural equation modeling.

Proposition Cluster 2 (Joint Optimization to Governance): Responsible AI Governance quality is predicted primarily by Joint Optimization quality, with additional direct contributions from Ethical Leadership (which establishes normative foundations) and Digital Leadership (which provides technical enablement). The mediating role of Joint Optimization will be tested empirically.

Proposition Cluster 3 (Governance to Outcomes): Organizational AI Outcomes measured through patient safety indicators, equity metrics, and workforce wellbeing are determined primarily by Responsible AI Governance quality and Joint Optimization, with the latter exerting both direct effects and indirect effects mediated through governance structures.

The specific path coefficients, effect sizes, and variance explained by each construct will be determined through quantitative analysis using Partial Least Squares Structural Equation Modeling (PLS-SEM) with the healthcare leader sample described in Section 3. **Table 4. Theoretical propositions of the Responsible AI Leadership Framework (Khompi's RALF).**

ID	Proposition
P1	Ethical leadership is positively associated with joint optimization quality in AI-integrated healthcare.
P2	Adaptive leadership is positively associated with joint optimization quality, especially where AI alters professional roles.
P3	Digital leadership is positively associated with joint optimization quality via enabling conditions for AI adoption.

P4	AI technology acceptance facilitation is positively associated with joint optimization quality.
P5	Joint optimization quality mediates the relationships between leadership constructs and responsible AI governance.
P6	Responsible AI governance quality is positively associated with organizational AI outcomes (safety, equity, wellbeing).
P7	Resource constraints moderate the JO–governance relationship; frugal settings require more intentional leadership.

I emphasize that Khompi’s RALF is presented here as a conceptual framework with testable theoretical propositions, not as a validated measurement model. The empirical validation including construct operationalization, measurement model assessment, and structural path estimation constitutes the next phase of this research program. The framework’s value at this stage lies in its theoretical integration of STS principles with leadership scholarship and its specification of testable relationships among constructs that have not previously been integrated in healthcare AI governance research.

Joint Optimization is determined by the combined influence of Ethical Leadership, Adaptive Leadership, Digital Leadership, and AI Technology Acceptance Facilitation. Responsible AI Governance quality is predicted by Joint Optimization quality, with additional direct contributions from Ethical and Digital Leadership. Organizational AI Outcomes are determined primarily by Responsible AI Governance quality and Joint Optimization

The framework will be empirically validated through structural equation modeling in future quantitative research, with specific path coefficients and effect sizes to be determined through data analysis.

4.6 Illustrative Application: Joint Optimization in Clinical AI Governance

To illustrate Khompi’s RALF’s applied implications, let’s explore the Epic Sepsis Model, deployed across hundreds of hospitals, was found by Wong et al. to have poor positive predictive value in external validation, illustrating precisely the leadership gap Khompi’s RALF addresses: the tool was widely deployed based on vendor claims without the joint-optimization governance that would have required independent equity auditing and clinician-informed evaluation before scaling [37]. A technical-optimization-only leadership approach would focus on algorithm selection (optimizing for sensitivity and specificity on validation data), integration with the electronic health record, and alert delivery in the clinical workflow. These are necessary but insufficient conditions for responsible deployment.

A Khompi’s RALF-informed approach would proceed differently. Ethical leadership would require equity auditing of the sepsis prediction model’s performance across patient demographic subgroups before deployment, establishing that the algorithm does not disadvantage specific populations. Adaptive leadership would convene ICU clinicians before deployment not merely to announce the new tool but to surface and process the professional anxieties it raises what does it mean for clinical expertise when an algorithm flags sepsis earlier than experienced nurses? Digital leadership would ensure adequate infrastructure (EHR integration, alert fatigue management, performance dashboards), appropriate implementation pace, and committed technical support. AI technology acceptance facilitation would address adoption through structured training, peer champion programs, and transparent communication of performance data after deployment. Joint optimization would be operationalized through a governance structure that includes clinical representation, regular equity reviews, patient feedback mechanisms, and workforce wellbeing monitoring alongside technical performance tracking.

This illustrative case is not merely pedagogical. It demonstrates that Khompi's RALF does not slow AI deployment for its own sake; it resequences and reweights the work of deployment so that social system conditions are treated as design requirements rather than afterthoughts. In frugal settings where a failed AI deployment consumes scarce capital, erodes already-thin clinician trust, and may worsen equity gaps this resequencing is not optional. It is the difference between AI that multiplies human capability and AI that multiplies organizational fragility.

5 Discussion

5.1 Implications for Healthcare Leadership Practice and Policy

If my framework is validated empirically, I expect to carry several significant implications for healthcare leadership practice. First, it reframes what "good" AI leadership looks like. Current organizational practice predominantly evaluates AI initiatives on technical and operational metrics: algorithm performance scores, implementation timelines, cost savings, and clinician throughput. Khompi's RALF implies that these metrics are necessary but insufficient, and that responsible AI leadership must be evaluated on social system outcomes as well: clinician trust and wellbeing, equity of algorithmic recommendations, quality of informed consent processes, and the degree to which clinician professional identity is protected rather than eroded by AI deployment [5, 6, 23, 24].

Second, Khompi's RALF implies specific leadership development priorities for healthcare organizations investing in AI. If joint optimization quality is the proximate determinant of responsible AI governance, then leadership development must explicitly build the competencies that drive joint optimization: ethical reasoning about algorithmic decisions, adaptive facilitation of professional identity transition, digital literacy adequate for sound AI governance decisions, and skills for facilitating technology acceptance without manufacturing consent. These are not competencies currently emphasized in most healthcare leadership development programs, which tend to focus on strategic planning, financial management, and quality improvement methodologies from the pre-AI era [22].

Third, Khompi's RALF has policy implications for healthcare AI regulation and accreditation. If responsible AI governance depends on leadership quality as much as algorithmic quality, regulatory frameworks that focus exclusively on algorithm testing and certification miss a significant determinant of deployment outcomes. Accreditation bodies and health system regulators might consider requiring organizations deploying clinical AI to demonstrate leadership governance structures as a condition of authorization, analogous to the requirement for institutional ethics committees as a condition of conducting human subjects research [24, 27]. Such a requirement would institutionalize joint optimization as a governance norm rather than leaving it to the discretionary virtue of individual leaders.

For example, the Mayo Clinic's Center for Digital Health has established an AI governance structure that includes clinician representatives, ethicists, and patient advocates in deployment decisions, a model consistent with Khompi's RALF's joint optimization requirement. Similarly, the Veterans Health Administration's National AI Institute has implemented phased deployment protocols that align with the adaptive leadership dimension of Khompi's RALF [39].

5.2 Limitations

Several limitations bear explicit acknowledgment. The proposed study will recruit participants from AI-deploying healthcare organizations in the United States, which may limit generalizability to health systems operating under different regulatory, cultural, and resource conditions. Future research should replicate and extend Khompi's RALF in international contexts, including low- and middle-income country health systems where AI deployment under resource constraint is most acute and where frugal AI governance frameworks are most urgently needed. The reliance on self-reported leadership behavior in the quantitative phase introduces social desirability bias, which is partially but not fully mitigated by the multi-informant qualitative phase. Objective organizational outcomes (patient safety indicators, equity metrics) from administrative data sources would strengthen future research designs. Finally, Khompi's RALF is proposed as a leadership governance framework, not a complete theory of healthcare AI outcomes; it is situated within a complex system of technological, regulatory, clinical, and organizational determinants that future research should model more comprehensively [20, 28].

6 Conclusion and Future Scope

The evidence presented in this paper suggests that leadership quality, not algorithmic sophistication, is the primary determinant of whether healthcare AI deployments succeed or fail. The technology is sufficiently mature to change the practice of medicine at scale; the evidence of what happens when it is deployed without adequate leadership and governance is equally mature. The 70–90% failure rate of digital health initiatives is not a technical verdict it is a leadership indictment. It tells us that the organizations deploying AI have not yet developed the leadership capabilities required to do so responsibly: to jointly optimize the human and technical systems on which effective, equitable, and trustworthy AI depends.

The Khompi's Responsible AI Leadership Framework presented in this paper addresses this gap by grounding healthcare AI governance in Socio-Technical Systems theory and operationalizing joint optimization through four empirically validated leadership constructs: ethical leadership, adaptive leadership, digital leadership, and AI technology acceptance facilitation. Khompi's RALF provides healthcare leaders with a theoretically grounded and practically actionable framework for governing AI deployment in ways that advance both technological capability and human values. It provides researchers with a testable structural model linking leadership behaviors to governance outcomes and, through them, to the patient safety, clinical equity, and workforce wellbeing indicators that constitute the ultimate measure of healthcare AI's societal value.

The stakes are not abstract. Every healthcare leader who deploys AI without attending to joint optimization makes a bet that algorithmic capability will substitute for organizational and cultural readiness a bet the evidence consistently shows is losing. Every leader who attends to both subsystems demonstrates what responsible AI governance looks like in practice: not slower AI, not more cautious AI, but human-centered AI, governed by leaders who understand that the measure of a technology is not what it can do but what it enables human beings to become and to give to each other. Healthcare organizations that invest in joint optimization, building leadership capacity alongside algorithmic capability, will be better positioned to deploy AI that improves care quality, reduces inequity, and sustains workforce wellbeing.

6.1 Future Scope

I envision several avenues for future research follow directly from the framework I have proposed. First, empirical validation of Khompi's RALF through the mixed-methods design specified in Section 3 will test the structural propositions and refine construct operationalization. Second, cross-national comparative studies should examine whether the leadership-to-governance pathways hold under divergent regulatory regimes, cultural norms of clinical authority, and resource endowments particularly in low- and middle-income country health systems where frugal AI is the default condition rather than a design choice. Third, longitudinal designs tracking AI deployments from selection through multi-year post-implementation would illuminate how joint optimization quality evolves over time and whether early leadership investments compound into durable governance advantages. Fourth, multi-level models incorporating team-level and organizational-level moderators (e.g., psychological safety climate, board-level AI oversight) would extend Khompi's RALF from a leader-centric to a systems-centric account of responsible AI. Fifth, intervention studies that train healthcare leaders in Khompi's RALF competencies and evaluate subsequent changes in governance quality and AI outcomes would establish the framework's practical utility as a leadership development curriculum.

The theoretical tools exist. The empirical pathway is specified. What is now needed is the organizational courage to lead AI deployment as the human enterprise it fundamentally is and the scholarly commitment to generate the evidence that makes such leadership not merely aspirational, but teachable, measurable, and scalable.

References

- [1] Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019). Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464), 447–453. <https://doi.org/10.1126/science.aax2342>
- [2] Topol, E. J. (2019). High-performance medicine: The convergence of human and artificial intelligence. *Nature Medicine*, 25(1), 44–56. <https://doi.org/10.1038/s41591-018-0300-7>
- [3] MarketsandMarkets. (2022). *Artificial intelligence in healthcare market worth \$67.4 billion by 2027: Exclusive report by MarketsandMarkets*
- [4] McKinsey & Company. (2018, October 29). *Unlocking success in digital transformations*. McKinsey & Company.
- [5] Brown, M. E., Treviño, L. K., & Harrison, D. A. (2005). Ethical leadership: A social learning perspective for construct development and testing. *Organizational Behavior and Human Decision Processes*, 97(2), 117–134. <https://doi.org/10.1016/j.obhdp.2005.03.002>
- [6] Heifetz, R. A., Grashow, A., & Linsky, M. (2009). *The practice of adaptive leadership: Tools and tactics for changing your organization and the world*. Harvard Business Press.
- [7] Zeike, S., Bradbury, K., Lindert, L., & Pfaff, H. (2019). Digital leadership skills and associations with psychological well-being. *International Journal of Environmental Research and Public Health*, 16(14), 2628. <https://doi.org/10.3390/ijerph16142628>
- [8] Trist, E. L., & Bamforth, K. W. (1951). Some social and psychological consequences of the longwall method of coal-getting. *Human Relations*, 4(1), 3–38. <https://doi.org/10.1177/001872675100400101>
- [9] Cherns, A. (1976). The principles of sociotechnical design. *Human Relations*, 29(8), 783–792. <https://doi.org/10.1177/001872677602900806>
- [10] Sittig, D. F., & Singh, H. (2010). A new sociotechnical model for studying health information technology in complex adaptive healthcare systems. *Quality and Safety in Health Care*, 19(Suppl 3), i68–i74. <https://doi.org/10.1136/qshc.2010.042085>

- [11] Rahimi, B., & Nadri, H. (2021). Health informatics leadership competencies: A systematic review. *Journal of Healthcare Leadership*, 13, 183–199. <https://doi.org/10.2147/JHL.S324082>
- [12] Venkatesh, V., Thong, J. Y. L., & Xu, X. (2012). Consumer acceptance and use of information technology: Extending the unified theory of acceptance and use of technology. *MIS Quarterly*, 36(1), 157–178. <https://doi.org/10.2307/41410412>
- [13] Shin, D. (2021). The effects of explainability and causability on perception, trust, and acceptance: Implications for explainable AI. *International Journal of Human-Computer Studies*, 146, 102551. <https://doi.org/10.1016/j.ijhcs.2020.102551>
- [14] Mittelstadt, B. D., & Floridi, L. (2016). The ethics of big data: Current and foreseeable issues in biomedical contexts. *Science and Engineering Ethics*, 22(2), 303–341. <https://doi.org/10.1007/s11948-015-9652-2>
- [15] Jiang, F., Jiang, Y., Zhi, H., Dong, Y., Li, H., Ma, S., Wang, Y., Dong, Q., Shen, H., & Wang, Y. (2017). Artificial intelligence in healthcare: Past, present and future. *Stroke and Vascular Neurology*, 2(4), 230–243. <https://doi.org/10.1136/svn-2017-000101>
- [16] Bates, D. W., Saria, S., Ohno-Machado, L., Shah, A., & Escobar, G. (2014). Big data in health care: Using analytics to identify and manage high-risk and high-cost patients. *Health Affairs*, 33(7), 1123–1131. <https://doi.org/10.1377/hlthaff.2014.0041>
- [17] Creswell, J. W., & Creswell, J. D. (2018). *Research design: Qualitative, quantitative, and mixed methods approaches* (5th ed.). SAGE Publications.
- [18] Hair, J. F., Risher, J. J., Sarstedt, M., & Ringle, C. M. (2019). When to use and how to report results of PLS-SEM. *European Business Review*, 31(1), 2–24. <https://doi.org/10.1108/EBR-11-2018-0203>
- [19] King, N. (2004). Using templates in the thematic analysis of text. In C. Cassell & G. Symon (Eds.), *Essential guide to qualitative methods in organizational research* (pp. 256–270). SAGE Publications.
- [20] Sittig, D. F., & Singh, H. (2015). A red-yellow-green patient safety taxonomy to identify and address health information technology–related patient safety concerns. *Journal of the American Medical Informatics Association*, 22(6), 1154–1160. <https://doi.org/10.1093/jamia/ocv048>
- [21] Lupton, D. (2014). The commodification of patient opinion: The digital patient experience economy in the age of big data. *Sociology of Health & Illness*, 36(6), 856–869. <https://doi.org/10.1111/1467-9566.12109>
- [22] Sinsky, C., Colligan, L., Li, L., Prgomet, M., Reynolds, S., Goeders, L., Westbrook, J., Tutty, M., & Blike, G. (2016). Allocation of physician time in ambulatory practice: A time and motion study in 4 specialties. *Annals of Internal Medicine*, 165(11), 753–760. <https://doi.org/10.7326/M16-0961>
- [23] Agrawal, A., Gans, J., & Goldfarb, A. (2018). *Prediction machines: The simple economics of artificial intelligence*. Harvard Business Review Press.
- [24] World Health Organization. (2021). *Ethics and governance of artificial intelligence for health: WHO guidance*. World Health Organization. <https://www.who.int/publications/i/item/9789240029200>
- [25] Jobin, A., Ienca, M., & Vayena, E. (2019). The global landscape of AI ethics guidelines. *Nature Machine Intelligence*, 1(9), 389–399. <https://doi.org/10.1038/s42256-019-0088-2>
- [26] Floridi, L., Cowls, J., Beltrametti, M., Chatila, R., Chazerand, P., Dignum, V., Luetge, C., Madelin, R., Pagallo, U., Rossi, F., Schafer, B., Valcke, P., & Vayena, E. (2018). AI4People—An ethical framework for a good AI society: Opportunities, risks, principles, and recommendations. *Minds and Machines*, 28(4), 689–707. <https://doi.org/10.1007/s11023-018-9482-5>
- [27] Dafoe, A. (2018). *AI governance: A research agenda*. Future of Humanity Institute, University of Oxford.
- [28] Collins, G. S., & Moons, K. G. M. (2019). Reporting of artificial intelligence prediction models. *The Lancet*, 393(10181), 1577–1579. [https://doi.org/10.1016/S0140-6736\(19\)30037-6](https://doi.org/10.1016/S0140-6736(19)30037-6)
- [29] Char et al., 2018, *New England Journal of Medicine*; or WHO, 2021.
- [30] Floridi, L., Cowls, J., Beltrametti, M., Chatila, R., Chazerand, P., Dignum, V., Luetge, C., Madelin, R., Pagallo, U., Rossi, F., Schafer, B., Valcke, P., & Vayena, E. (2018). AI4People—An ethical

framework for a good AI society: Opportunities, risks, principles, and recommendations. *Mind and Machines*, 28(4), 689–707. <https://doi.org>

- [31] Cherns, A. (1976). The principles of sociotechnical design. *Human Relations*, 29(8), 783–792. <https://doi.org>
- [32] Trist, E. L., & Bamforth, K. W. (1951). Some social and psychological consequences of the longwall method of coal-getting: An examination of the psychological situation and defences of a work group in relation to the social structure and technological content of the work system. *Human Relations*, 4(1), 3–38. <https://doi.org>
- [33] Brown, M. E., Treviño, L. K., & Harrison, D. A. (2005). Ethical leadership: A social learning perspective for construct development and testing. *Organizational Behavior and Human Decision Processes*, 97(2), 117–134. <https://doi.org>
- [34] Heifetz, R. A., Heifetz, R., Grashow, A., & Linsky, M. (2009). *The practice of adaptive leadership: Tools and tactics for changing your organization and the world*. Harvard Business Press.
- [35] Venkatesh, V., Thong, J. Y., & Xu, X. (2012). Consumer acceptance and use of information technology: Extending the unified theory of acceptance and use of technology. *MIS Quarterly*, 36(1), 157–178. <https://doi.org>
- [36] Zeike, S., Bradbury, K., Lindert, L., & Pfaff, H. (2019). Digital leadership skills and associations with psychological well-being. *International Journal of Environmental Research and Public Health*, 16(14), Article 2628. <https://doi.org>
- [37] Wong, A., Otles, E., Donnelly, J. P., Krumm, A., McCullough, J. S., DeTroyer-Cooley, O., Pestru, J., Phillips, M., Konye, A., Athey, M., & Wiens, J. (2021). External validation of a widely implemented proprietary sepsis prediction model in hospitalized patients. *JAMA Internal Medicine*, 181(8), 1065–1074. <https://doi.org/10.1001/jamainternmed.2021.2626>
- [38] Larrazabal, A. J., Nieto, N., Peterson, V., Milone, D. H., & Ferrante, E. (2020). Gender imbalance in medical imaging datasets produces biased classifiers for computer-aided diagnosis. *Proceedings of the National Academy of Sciences*, 117(23), 12592–12594. <https://doi.org/10.1073/pnas.1919012117>
- [39] Mayo Clinic. (2022). *Center for Digital Health and AI governance framework*. Mayo Clinic Health System. <https://www.mayoclinic.org/about-mayo-clinic/governance>